

Addison Kalanther (addikala@berkeley.edu)

## Lecture 3: PyTorch

This is a quick overview of PyTorch, an autodiff framework used to train models.

### What is a model?

A model is of the form

$$\theta^* = \operatorname{argmin}_{\theta} \sum_{(x,y) \in D} \mathcal{L}(f_{\theta}(x), y)$$

with

- $\operatorname{argmin}_{\theta}$  being the gradient descent,
- $D$  being the dataset,
- $\mathcal{L}$  being the loss function, and
- $f_{\theta}$  being the neural network.

PyTorch does all of these and more.

### Using GPU

NumPy is CPU-only and does not support autodiff. PyTorch allows for GPU use to parallelize processes and supports autodiff.

To move a PyTorch object to GPU, use `.to("cuda")` / `.cuda()`.

To move a PyTorch object to CPU, use `.to("cpu")` / `.cpu()`.

(New) To move a PyTorch object to MPS (Metal backend for Apple ARM processors), use `.to(torch.device("mps"))`.

### Computing gradients

To compute the gradient of a Tensor, ensure `requires_grad=True`. Now, we can find the gradient of these variables when taking a gradient with respect to another value, most commonly loss.

Graphs for calculating gradients take up a lot of memory. To detach gradient values from Tensor, use `.detach()`.

### Training loop

If there is anything to remember in PyTorch, it is these 3 lines.

```
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

`optimizer.zero_grad()` resets the optimizer to zero, `loss.backward()` computes the gradients with respect to loss, and `optimizer.step()` takes a step in the direction opposite of the gradients calculated by the `.backward()` call to perform gradient descent.

### Example training loop

```
for epoch in range(num_epochs):
    net.train() # puts net in training mode
    for data, target in dataloader:
        data = torch.from_numpy(data).float().cuda()
        target = torch.from_numpy(target).float().cuda()

        prediction = net(data)
        loss = loss_fn(prediction, target) # loss takes (pred, truth)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    net.eval()
    # Do evaluation...
```

### NumPy / PyTorch conversion

Numpy -> PyTorch: `torch.from_numpy(numpy_array).float()`

- `.float()` if trying to convert to float, most common for neural nets

PyTorch -> NumPy: `torch_tensor.detach().cpu().numpy()`

- `.detach()` only necessary if `require_grad=True`
- `.cpu()` only necessary if not on CPU